

grebena_regresija

August 23, 2017

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn import linear_model
from sklearn import metrics
from sklearn import model_selection
%matplotlib inline
```

```
In [2]: # radimo sa podacima koji se ticu bejzbol igraca
# hteli bismo na osnovu svih statistika o performansama igraca da predvidimo
# podaci koji se koriste su preuzeti iz knjige Introduction to Statistical
# link do knjige je: http://www-bcf.usc.edu/~gareth/ISL/
data = pd.read_csv('hitters.csv')
```

```
In [3]: # skup sadrzi 21 atribut (ukljucujuci i platu koju zelimo da predvidimo)
# i 322 informacije koje se ticu pojedinacnih igraca
data.shape
```

Out[3]: (322, 21)

```
In [4]: data.head(5)
```

```
Out[4]:
```

	Player	AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAtBat
0	-Andy Allanson	293	66	1	30	29	14	1	293
1	-Alan Ashby	315	81	7	24	38	39	14	3449
2	-Alvin Davis	479	130	18	66	72	76	3	1624
3	-Andre Dawson	496	141	20	65	78	37	11	5628
4	-Andres Galarraga	321	87	10	39	42	30	2	396

	CHits	...	CRuns	CRBI	CWalks	League	Division	PutOuts	Assists
0	66	...	30	29	14	A	E	446	33
1	835	...	321	414	375	N	W	632	43
2	457	...	224	266	263	A	W	880	82
3	1575	...	828	838	354	N	E	200	11
4	101	...	48	46	33	N	E	805	40

	Errors	Salary	NewLeague
0	20	NaN	A

1	10	475.0	N
2	14	480.0	A
3	3	500.0	N
4	4	91.5	N

[5 rows x 21 columns]

```
In [5]: # ispitujemo prirodu podataka koji se nalaze u tabeli
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 322 entries, 0 to 321
Data columns (total 21 columns):
Player      322 non-null object
AtBat       322 non-null int64
Hits        322 non-null int64
HmRun       322 non-null int64
Runs        322 non-null int64
RBI         322 non-null int64
Walks       322 non-null int64
Years       322 non-null int64
CAtBat      322 non-null int64
CHits       322 non-null int64
CHmRun      322 non-null int64
CRuns       322 non-null int64
CRBI        322 non-null int64
CWalks      322 non-null int64
League      322 non-null object
Division    322 non-null object
PutOuts     322 non-null int64
Assists     322 non-null int64
Errors      322 non-null int64
Salary      263 non-null float64
NewLeague   322 non-null object
dtypes: float64(1), int64(16), object(4)
memory usage: 52.9+ KB
```

```
In [6]: # neke od vrsta zadrze nedostajuće vrednosti tako da nam nisu podesne u modelu
data = data.dropna()
```

```
In [7]: data.shape
```

```
Out[7]: (263, 21)
```

```
In [8]: # informacije o platama su nam potrebne za predviđanje kao posebna veličina
Y = data['Salary']
```

```
In [10]: # neki atributi nam nisu informativni - kategorickog su tipa, a ne numerickog
# to su Player, League, Division, NewLeague
```

```
# Salary takodje treba ukloniti
X = data.drop(['Player', 'League', 'Division', 'NewLeague', 'Salary'], axis=1)
print(X)
```

	AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAtBat	CHits	CHmRun	\
1	315	81	7	24	38	39	14	3449	835	69	
2	479	130	18	66	72	76	3	1624	457	63	
3	496	141	20	65	78	37	11	5628	1575	225	
4	321	87	10	39	42	30	2	396	101	12	
5	594	169	4	74	51	35	11	4408	1133	19	
6	185	37	1	23	8	21	2	214	42	1	
7	298	73	0	24	24	7	3	509	108	0	
8	323	81	6	26	32	8	2	341	86	6	
9	401	92	17	49	66	65	13	5206	1332	253	
10	574	159	21	107	75	59	10	4631	1300	90	
11	202	53	4	31	26	27	9	1876	467	15	
12	418	113	13	48	61	47	4	1512	392	41	
13	239	60	0	30	11	22	6	1941	510	4	
14	196	43	7	29	27	30	13	3231	825	36	
16	568	158	20	89	75	73	15	8068	2273	177	
17	190	46	2	24	8	15	5	479	102	5	
19	127	32	8	16	22	14	8	727	180	24	
20	413	92	16	72	48	65	1	413	92	16	
21	426	109	3	55	43	62	1	426	109	3	
23	472	116	16	60	62	74	6	1924	489	67	
24	629	168	18	73	102	40	18	8424	2464	164	
25	587	163	4	92	51	70	6	2695	747	17	
26	324	73	4	32	18	22	7	1931	491	13	
27	474	129	10	50	56	40	10	2331	604	61	
28	550	152	6	92	37	81	5	2308	633	32	
29	513	137	20	90	95	90	14	5201	1382	166	
31	419	108	6	55	36	22	3	591	149	8	
33	583	168	17	83	80	56	5	1646	452	44	
34	204	49	6	23	25	12	7	1309	308	27	
35	379	106	10	38	60	30	14	6207	1906	146	
..	
287	687	213	10	91	65	27	4	1518	448	15	
288	368	103	3	48	28	54	8	1897	493	9	
289	263	70	1	26	23	30	4	888	220	9	
290	642	211	14	107	59	52	5	2364	770	27	
291	265	68	8	26	30	29	7	1337	339	32	
293	559	141	2	48	61	73	8	3162	874	16	
294	520	120	17	53	44	21	4	927	227	22	
295	19	4	1	2	3	1	1	19	4	1	
296	205	43	2	24	17	20	7	854	219	12	
297	193	47	10	21	29	24	6	1136	256	42	
299	213	61	4	17	22	3	17	4061	1145	83	
300	510	147	10	56	52	53	7	2872	821	63	

301	578	138	1	56	59	34	3	1399	357	7
303	441	113	5	76	52	76	5	1546	397	17
304	172	42	3	17	14	15	10	4086	1150	57
306	127	32	4	14	25	12	19	8396	2402	242
307	279	69	4	35	31	32	4	1359	355	31
308	480	112	18	50	71	44	7	3031	771	110
309	600	139	0	94	29	60	2	1236	309	1
310	610	186	19	107	98	74	6	2728	753	69
311	360	81	5	37	44	37	7	2268	566	41
312	387	124	1	67	27	36	7	1775	506	6
313	580	207	8	107	71	105	5	2778	978	32
314	408	117	11	66	41	34	1	408	117	11
315	593	172	22	82	100	57	1	593	172	22
317	497	127	7	65	48	37	5	2703	806	32
318	492	136	5	76	50	94	12	5511	1511	39
319	475	126	3	61	43	52	6	1700	433	7
320	573	144	9	85	60	78	8	3198	857	97
321	631	170	9	77	44	31	11	4908	1457	30

	CRuns	CRBI	CWalks	PutOuts	Assists	Errors
1	321	414	375	632	43	10
2	224	266	263	880	82	14
3	828	838	354	200	11	3
4	48	46	33	805	40	4
5	501	336	194	282	421	25
6	30	9	24	76	127	7
7	41	37	12	121	283	9
8	32	34	8	143	290	19
9	784	890	866	0	0	0
10	702	504	488	238	445	22
11	192	186	161	304	45	11
12	205	204	203	211	11	7
13	309	103	207	121	151	6
14	376	290	238	80	45	8
16	1045	993	732	105	290	10
17	65	23	39	102	177	16
19	67	82	56	202	22	2
20	72	48	65	280	9	5
21	55	43	62	361	22	2
23	242	251	240	518	55	3
24	1008	1072	402	1067	157	14
25	442	198	317	434	9	3
26	291	108	180	222	3	3
27	246	327	166	732	83	13
28	349	182	308	262	329	16
29	763	734	784	267	5	3
31	80	46	31	226	7	4
33	219	208	136	109	292	25

34	126	132	66	419	46	5
35	859	803	571	72	170	24
..
287	196	137	89	294	445	13
288	207	162	198	209	246	3
289	83	82	86	81	147	4
290	352	230	193	337	19	4
291	135	163	128	92	5	3
293	421	349	359	352	414	9
294	106	80	52	70	144	11
295	2	3	1	692	70	8
296	105	99	71	131	6	1
297	129	139	106	299	13	5
299	488	491	244	178	45	4
300	307	340	174	810	99	18
301	149	161	87	133	371	20
303	226	149	191	160	290	11
304	579	363	406	65	0	0
306	1048	1348	819	167	18	6
307	180	148	158	133	173	9
308	338	406	239	94	270	16
309	201	69	110	300	12	9
310	399	366	286	1182	96	13
311	279	257	246	170	284	3
312	272	125	194	186	290	17
313	474	322	417	121	267	19
314	66	41	34	942	72	11
315	82	100	57	1222	139	15
317	379	311	138	325	9	3
318	897	451	875	313	381	20
319	217	93	146	37	113	7
320	470	420	332	1314	131	12
321	775	357	249	408	4	3

[263 rows x 16 columns]

In [13]: X.shape

Out[13]: (263, 16)

In [11]: # *posto su skale vrednosti podataka neujednacene, ima smisla raditi normalizaciju*

In [12]: # *delimo skup podataka na:*
training skup
validacioni skup koji koristimo za učenje metaparametara
test skup na kojem ocenjujemo uspesnost modela

```
X_train_validation, X_test, Y_train_validation, Y_test = model_selection.train_test_split(X_train, X_validation, Y_train, Y_validation, random_state=42)
```

```
In [13]: # baseline - linearna regresija
linreg = linear_model.LinearRegression(normalize=True)
linreg.fit(X_train_validation, Y_train_validation)
Y_predicted = linreg.predict(X_test)
linreg_score = metrics.mean_squared_error(Y_test, Y_predicted)
linreg_score
```

```
Out[13]: 138551.12033732756
```

```
In [14]: # grebena regresija sa metaparametrom 4
reg_1 = linear_model.Ridge(alpha=4, normalize=True)
reg_1.fit(X_train_validation, Y_train_validation)
Y_predicted = reg_1.predict(X_test)
reg_score_1 = metrics.mean_squared_error(Y_test, Y_predicted)
reg_score_1
# na osnovu ovog rezultata zakljucujemo da ima smisla raditi regularizaciju
```

```
Out[14]: 126928.00669545057
```

```
In [15]: # grebena regresija sa metaparametrom 10^10
reg_2 = linear_model.Ridge(alpha=10**10, normalize=True)
reg_2.fit(X_train_validation, Y_train_validation)
Y_predicted = reg_2.predict(X_test)
reg_score_2 = metrics.mean_squared_error(Y_test, Y_predicted)
reg_score_2
# na osnovu ovog rezultata zakljucujemo da ima smisla odrediti i precizno
# jer ocigledno nisu sve regularizacije dobre
```

```
Out[15]: 194839.49321934962
```

```
In [45]: # skup vrednosti iz kojeg trazimo najbolju vrednost za parametar alfa
# imamo i male i velike vrednosti u skupu
alphas = 10**np.linspace(10, -2, 100)*0.5
```

```
In [46]: alphas
```

```
Out[46]: array([[ 5.00000000e+09,  3.78231664e+09,  2.86118383e+09,
  2.16438064e+09,  1.63727458e+09,  1.23853818e+09,
  9.36908711e+08,  7.08737081e+08,  5.36133611e+08,
  4.05565415e+08,  3.06795364e+08,  2.32079442e+08,
  1.75559587e+08,  1.32804389e+08,  1.00461650e+08,
  7.59955541e+07,  5.74878498e+07,  4.34874501e+07,
  3.28966612e+07,  2.48851178e+07,  1.88246790e+07,
  1.42401793e+07,  1.07721735e+07,  8.14875417e+06,
  6.16423370e+06,  4.66301673e+06,  3.52740116e+06,
  2.66834962e+06,  2.01850863e+06,  1.52692775e+06,
```

```

1.15506485e+06, 8.73764200e+05, 6.60970574e+05,
5.00000000e+05, 3.78231664e+05, 2.86118383e+05,
2.16438064e+05, 1.63727458e+05, 1.23853818e+05,
9.36908711e+04, 7.08737081e+04, 5.36133611e+04,
4.05565415e+04, 3.06795364e+04, 2.32079442e+04,
1.75559587e+04, 1.32804389e+04, 1.00461650e+04,
7.59955541e+03, 5.74878498e+03, 4.34874501e+03,
3.28966612e+03, 2.48851178e+03, 1.88246790e+03,
1.42401793e+03, 1.07721735e+03, 8.14875417e+02,
6.16423370e+02, 4.66301673e+02, 3.52740116e+02,
2.66834962e+02, 2.01850863e+02, 1.52692775e+02,
1.15506485e+02, 8.73764200e+01, 6.60970574e+01,
5.00000000e+01, 3.78231664e+01, 2.86118383e+01,
2.16438064e+01, 1.63727458e+01, 1.23853818e+01,
9.36908711e+00, 7.08737081e+00, 5.36133611e+00,
4.05565415e+00, 3.06795364e+00, 2.32079442e+00,
1.75559587e+00, 1.32804389e+00, 1.00461650e+00,
7.59955541e-01, 5.74878498e-01, 4.34874501e-01,
3.28966612e-01, 2.48851178e-01, 1.88246790e-01,
1.42401793e-01, 1.07721735e-01, 8.14875417e-02,
6.16423370e-02, 4.66301673e-02, 3.52740116e-02,
2.66834962e-02, 2.01850863e-02, 1.52692775e-02,
1.15506485e-02, 8.73764200e-03, 6.60970574e-03,
5.00000000e-03])

```

```

In [47]: # niz gresaka nasih modela
errors = []

```

```


In [48]: # izracunavamo gresku za skvaku vrednost metaparametra alfa
for a in alphas:
    reg = linear_model.Ridge(normalize=True, alpha=a)
    reg.fit(X_train, Y_train)
    Y_predicted = reg.predict(X_validation)
    error = metrics.mean_squared_error(Y_predicted, Y_validation)
    errors.append(error)

```

```

In [49]: # iscrtavamo zavisnost greske od vrednosti metaparametra alfa
plt.plot(alphas, errors)
plt.xscale('log')

```



```
In [50]: # trazimo koja je to vrednost
```

```
In [51]: type(errors)
```

```
Out[51]: list
```

```
In [52]: erros = np.array(errors)
```

```
In [53]: erros.argmin()
```

```
Out[53]: 74
```

```
In [54]: alpha = alphas[74]
         alpha
```

```
Out[54]: 5.3613361100516048
```

```
In [55]: # greska modela
         errors[74]
```

```
Out[55]: 124988.0740587797
```

```
In [56]: # sada kada znamo najbolju vrednost metaparametra alfa mozemo da naucimo i
         # da damo konacnu procenu greske
         # ne zaboraviti da se u ovom koraku koriste i trening i validacioni skup

         regfin = linear_model.Ridge(alpha=5.361, normalize=True)
```



```
regfin.fit(X_train_validation, Y_train_validation)
Y_predicted = regfin.predict(X_test)
regfin_score = metrics.mean_squared_error(Y_test, Y_predicted)
regfin_score
```

Out[56]: 131350.87475923024

In []: